



Federated Averaging

Outline

- Motivation
- Main Works
 1. FederatedAveraging
 2. FedAvg with shared data
 3. LoAdaBoost FedAvg
 4. Asynchronous Federated Optimization
- Conclusion



Motivation

key properties different from a typical distributed optimization problem:

- Non-IID
- Unbalanced
- Massively Distributed
- Limited Communication



Training process




Step 1: participants locally compute training gradients, mask a selection of gradients with encryption , differential privacy or secret sharing techniques, and send masked results to server

Step 2: Server performs secure aggregation without learning information about any participant

Step 3: Server send back the aggregated results to participants

Step 4: Participants update their respective model with the decrypted gradients



The FederatedAveraging Algorithm



Algorithm 1 FederatedAveraging. The K clients are indexed by k ; B is the local minibatch size, E is the number of local epochs, and η is the learning rate.

Server executes:

```
initialize  $w_0$ 
for each round  $t = 1, 2, \dots$  do
   $m \leftarrow \max(C \cdot K, 1)$ 
   $S_t \leftarrow$  (random set of  $m$  clients)
  for each client  $k \in S_t$  in parallel do
     $w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$ 
   $w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$ 
```

```
ClientUpdate( $k, w$ ): // Run on client  $k$ 
   $\mathcal{B} \leftarrow$  (split  $\mathcal{P}_k$  into batches of size  $B$ )
  for each local epoch  $i$  from 1 to  $E$  do
    for batch  $b \in \mathcal{B}$  do
       $w \leftarrow w - \eta \nabla \ell(w; b)$ 
  return  $w$  to server
```

- Update (central server)

$$w_{t+1} \leftarrow w_t - \eta \sum_{k=1}^K \frac{n_k}{n} g_k,$$

- Update (client)

$$\forall k, w_{t+1}^k \leftarrow w_t - \eta g_k$$

- Three key parameters
 - **C** (the fraction of clients that perform computation on each round)
 - **E** (the number of training passes each client makes over its local dataset on each round)
 - **B** (the local minibatch size used for the client updates)

FedAvg with shared data

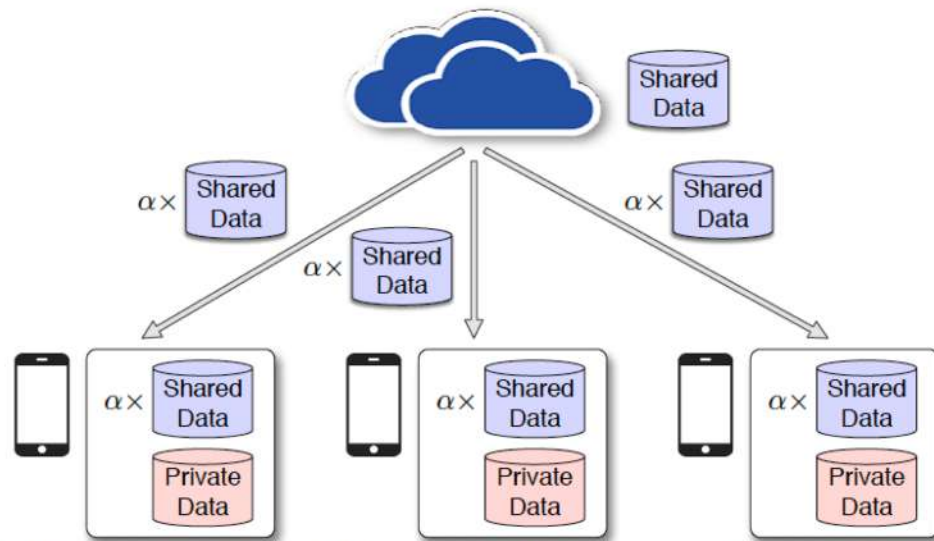


Figure 6: Illustration of the data-sharing strategy.

- A globally shared dataset G
- The warm-up model trained on G
- A random α portion of G are distributed to each client
- Each client is trained on the shared data together with the private data

FedAvg with shared data

$$\beta = \frac{\|G\|}{\|D\|} \times 100\%$$




Figure 7: (a) Test accuracy and EMD vs. β (b) Test accuracy vs. the distributed fraction α


- The test accuracy increases as β increases. Even with a lower $\beta = 10\%$, we can still achieve a test accuracy.
- Only a random portion of G needs to be distributed to each client.

LoAdaBoost FedAvg



- Loss-based Adaptive Boosting Federated Averaging
- Focus on medical data
- Three issues in federation learning:
 1. the local client-side computation complexity
 2. the communication cost
 3. the test accuracy

LoAdaBoost FedAvg



- Different from FedAvg, the learning process average was performed for $E/2$ instead of E epochs.

- $\Delta L_k^{t,0} = L_k^{t,0} - L_{median}^{t-1} \leq 0$ finish
- Otherwise, retrain for $E/2-1$ more epochs, (repeated for $E/2-r+1$ epochs), stop until
- $\Delta L_k^r = L_k^{t,r} - L_{median}^{t-1} \leq 0$ or total epochs reached $3E/2$


Asynchronous Federated Optimization



- Previous: Synchronous
 1. Too many devices checking in at the same time
 2. Difficult to synchronize the selected devices at the end of each epoch
- Asynchronous federated optimization
The key idea : use a weighted average to update the global model.

$$x_t = (1 - \alpha)x_{t-1} + \alpha x_{new},$$

Asynchronous Federated Optimization



Algorithm 1 Asynchronous Federated Optimization (FedAsync)

Server Process

Input: $\alpha \in (0, 1)$

Initialize $x_0, \alpha_t \leftarrow \alpha, \forall t \in [T]$

Scheduler Thread

Scheduler periodically triggers some training tasks on some workers, and sends them the latest global model with time stamp

Updater Thread

for all epoch $t \in [T]$ **do**

Receive the pair (x_{new}, τ) from any worker

Optional: $\alpha_t \leftarrow \alpha \times s(t - \tau)$, $s(\cdot)$ is a function of the staleness

$$x_t = (1 - \alpha_t)x_{t-1} + \alpha_t x_{new}$$

end for

Worker Processes

for all $i \in [n]$ in parallel **do**

If triggered by the scheduler:

Receive the pair of the global model and its time stamp (x_t, t) from the server

$\tau \leftarrow t, x_{\tau,0}^i \leftarrow x_t$

For μ -weakly convex F :

Define $g_{x_t}(x; z) = f(x; z) + \frac{\rho}{2}\|x - x_t\|^2$, where $\rho > \mu$

for all local iteration $h \in [H_\tau^i]$ **do**

Randomly sample $z_{\tau,h}^i \sim \mathcal{D}^i$

$$x_{\tau,h}^i \leftarrow \begin{cases} \text{Option I, for strongly convex } F: \\ x_{\tau,h-1}^i - \gamma \nabla f(x_{\tau,h-1}^i; z_{\tau,h}^i) \\ \text{Option II, for weakly convex } F: \\ x_{\tau,h-1}^i - \gamma \nabla g_{x_t}(x_{\tau,h-1}^i; z_{\tau,h}^i) \end{cases}$$

end for

Push $(x_{\tau,H_\tau^i}^i, \tau)$ to the server

end for