Vision Language Navigation

Weiwen Chen 2019.12.29

Outline

- Short history of VLN(from sim & data)
- Problem Definition
- Methods
  - Baseline: seq-to-seq
  - Look Before You Leap (model-based & model-free)
  - Speaker-Follower (Data Augmentation, Action Space)
  - Reinforced Cross-Modal Matching (Align matching, SIL)
  - Self-Monitoring (Align progress)
  - Environmental Dropout (Data Augmentation, IL)
  - Auxiliary Reasoning Tasks :)

History of VLN

Simulator & data

Follow Navigational Directions

1. go vertically down until you're underneath eh diamond mine
2. then eh go right until you're
3. you're between springbok and highest viewpoint

Figure 2: The instruction giver and instruction follower face each other, and cannot see each others maps.

Vogel, et al. Learning to Follow Navigational Directions. ACL2010

Mobile Manipulation

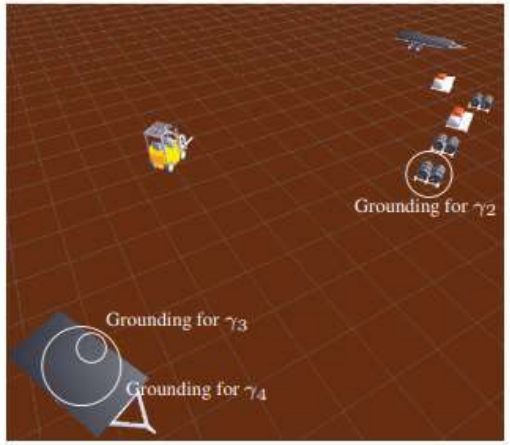

(a) Object groundings

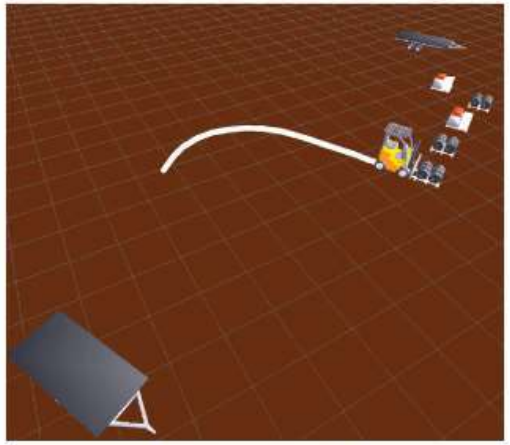
(b) Pick up the pallet

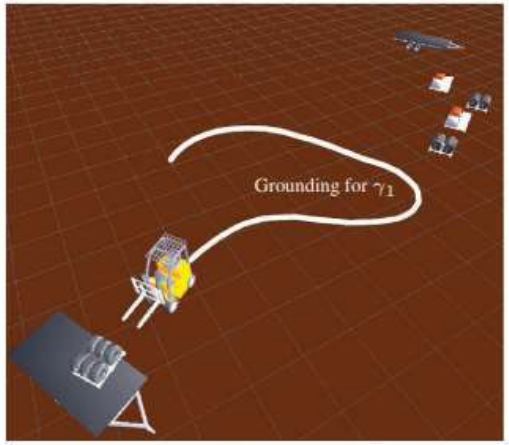

(c) Put it on the truck

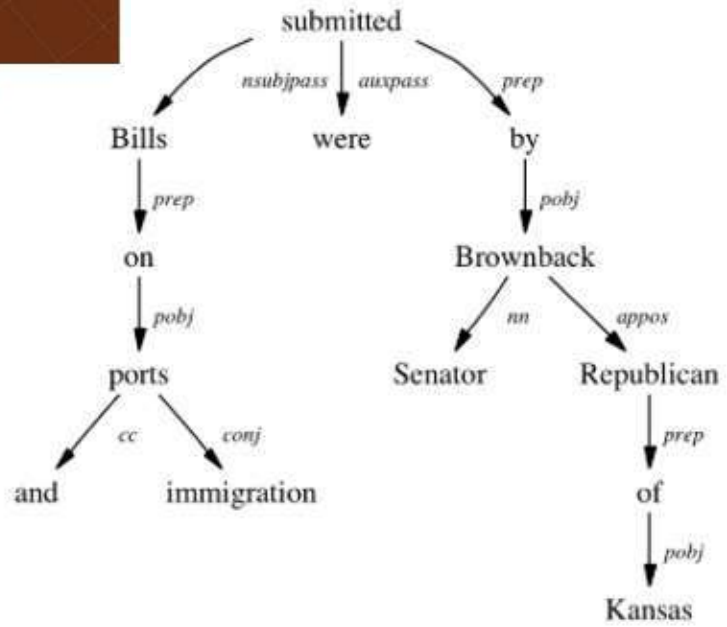

Stefanie Tellex, et al. Understanding Natural Language Commands for Robotic Navigation and Mobile Manipulation. AAAI2011

Language Navigation

Instruction: "Go away from the lamp to the intersection of the red brick and wood"

Basic:
Turn ( ) ,
Travel (steps: 1)

Landmarks:
Turn ( ) ,
Verify (left: WALL , back: LAMP , back: HATRACK , front: BRICK HALL) ,
Travel (steps: 1) ,
Verify (side: WOOD HALL)

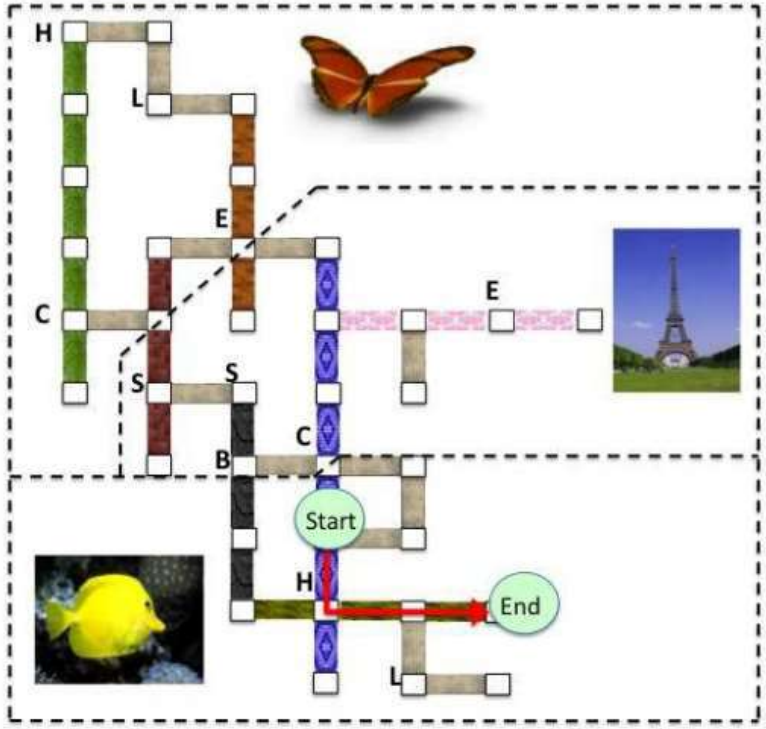

- "Go towards the coat rack and take a left at the coat rack. go all the way to the end of the hall and this is 4."
- "turn so that the wall is on your right side. walk forward once. turn left. walk forward twice."

Chen, David L, et al. Learning to interpret natural language navigation instructions from observations. AAAI2011

Gated-Attention

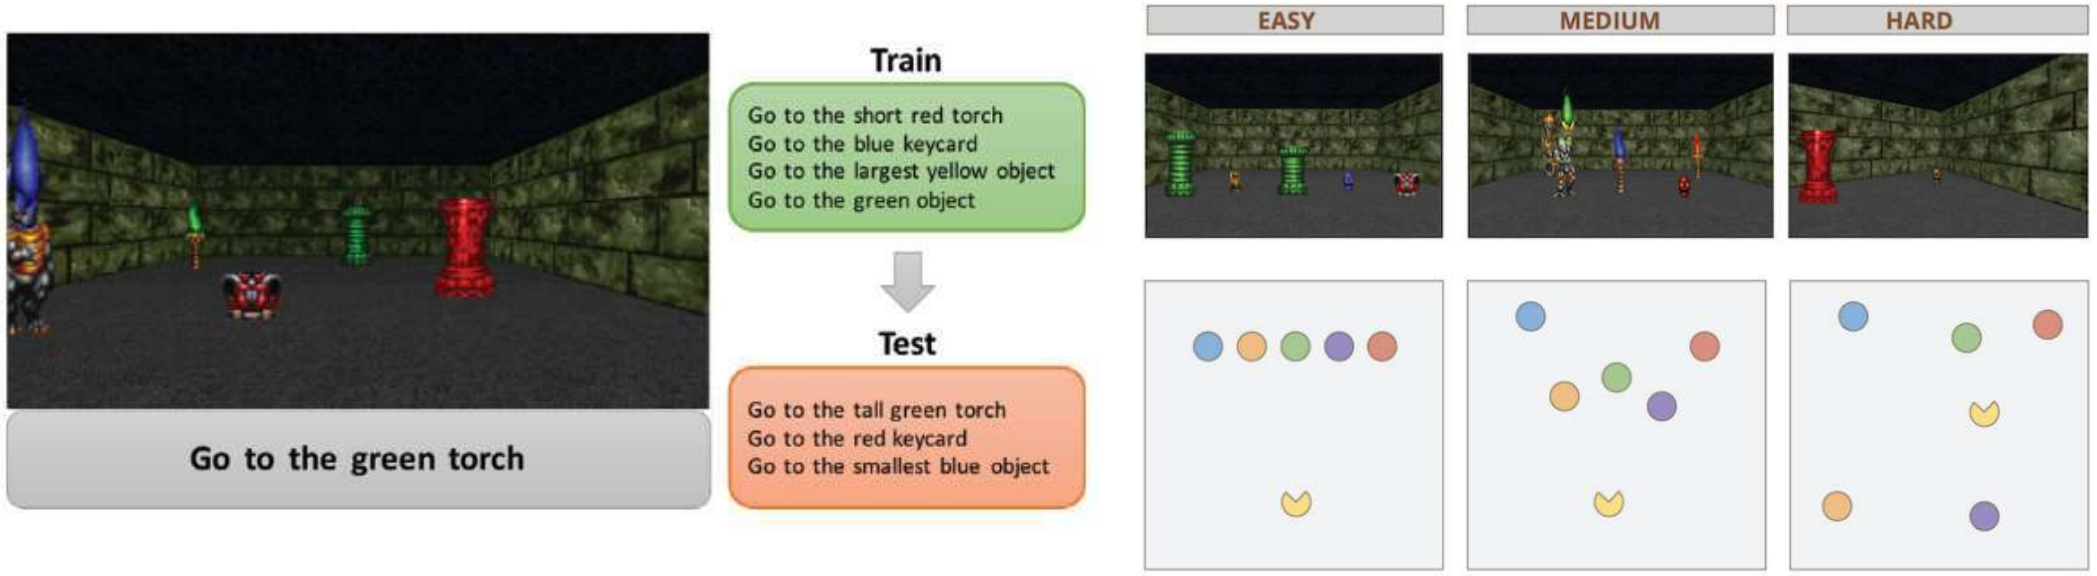

Chaplot, et al. Gated-Attention Architectures for Task-Oriented Language Grounding. AAAI2018
Sinha, et al. Attention Based Natural Language Grounding by Navigating Virtual Environment. WACV2019

Problem Def

costly

Matterport3D

- Big:
- 10,800 panoramic views
- from 194,400 RGB-D images
- Of 90 building-scale scenes

Chang, A,et al. Matterport3D: Learning from RGB-D data in indoor environments. 3DV 2017.

Vision-and-Language Navigation

- Over 400 workers
- 1,600 hours

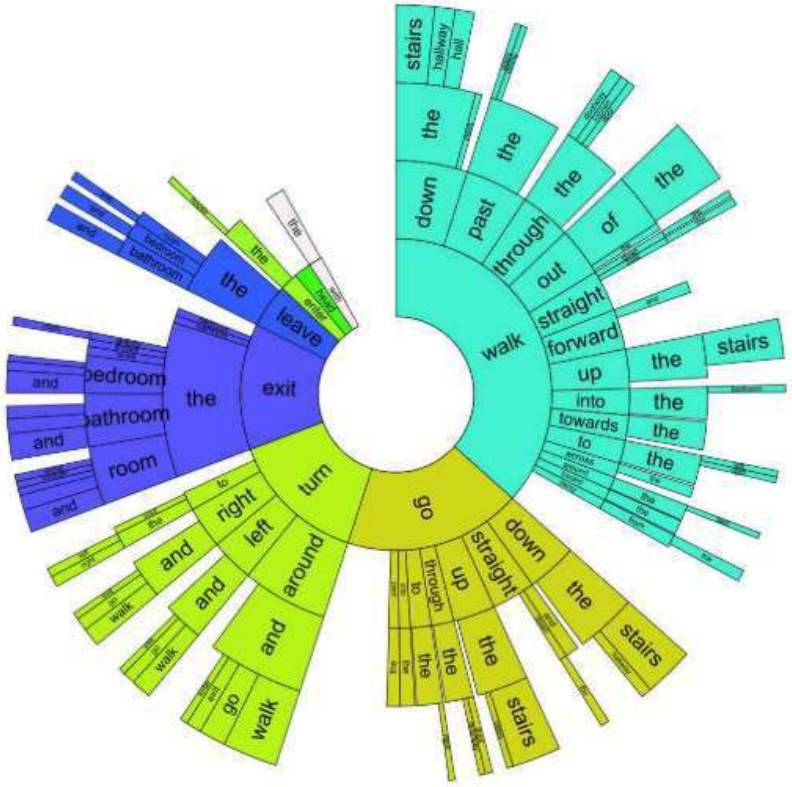

Instruction: Head upstairs and walk past the piano through an archway directly in front. Turn right when the hallway ends at pictures and table. Wait by the moose antlers hanging on the wall.

Anderson, et al. Vision-and-Language Navigation: Interpreting Visually-Grounded Navigation Instructions in Real Environments. CVPR2018

Action Space

- 6 actions
- View: left, right, up, down
- Move: forward
- End: stop

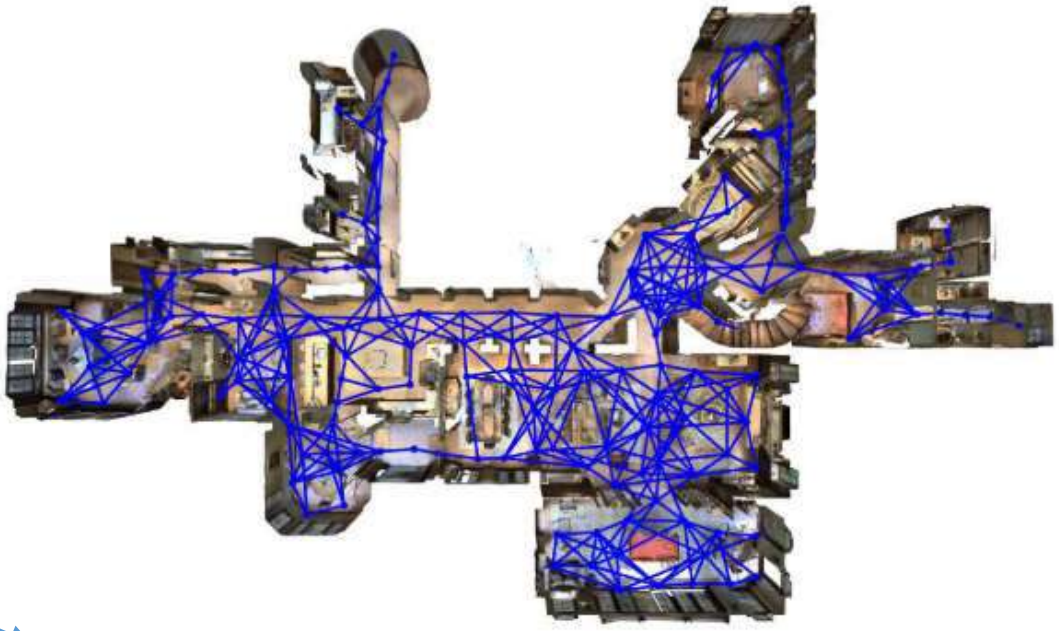

$$W_{t+1} = \{v_t\} \cup \{v_i \in V \mid \langle v_t, v_i \rangle \in E \wedge v_i \in P_t\}$$

Anderson, et al. Vision-and-Language Navigation: Interpreting Visually-Grounded Navigation Instructions in Real Environments. CVPR2018

Challenges

- cross-modal grounding
- ill-posed feedback
- Generalization

Methods

CV & NLP & RL

Baseline seq-to-seq

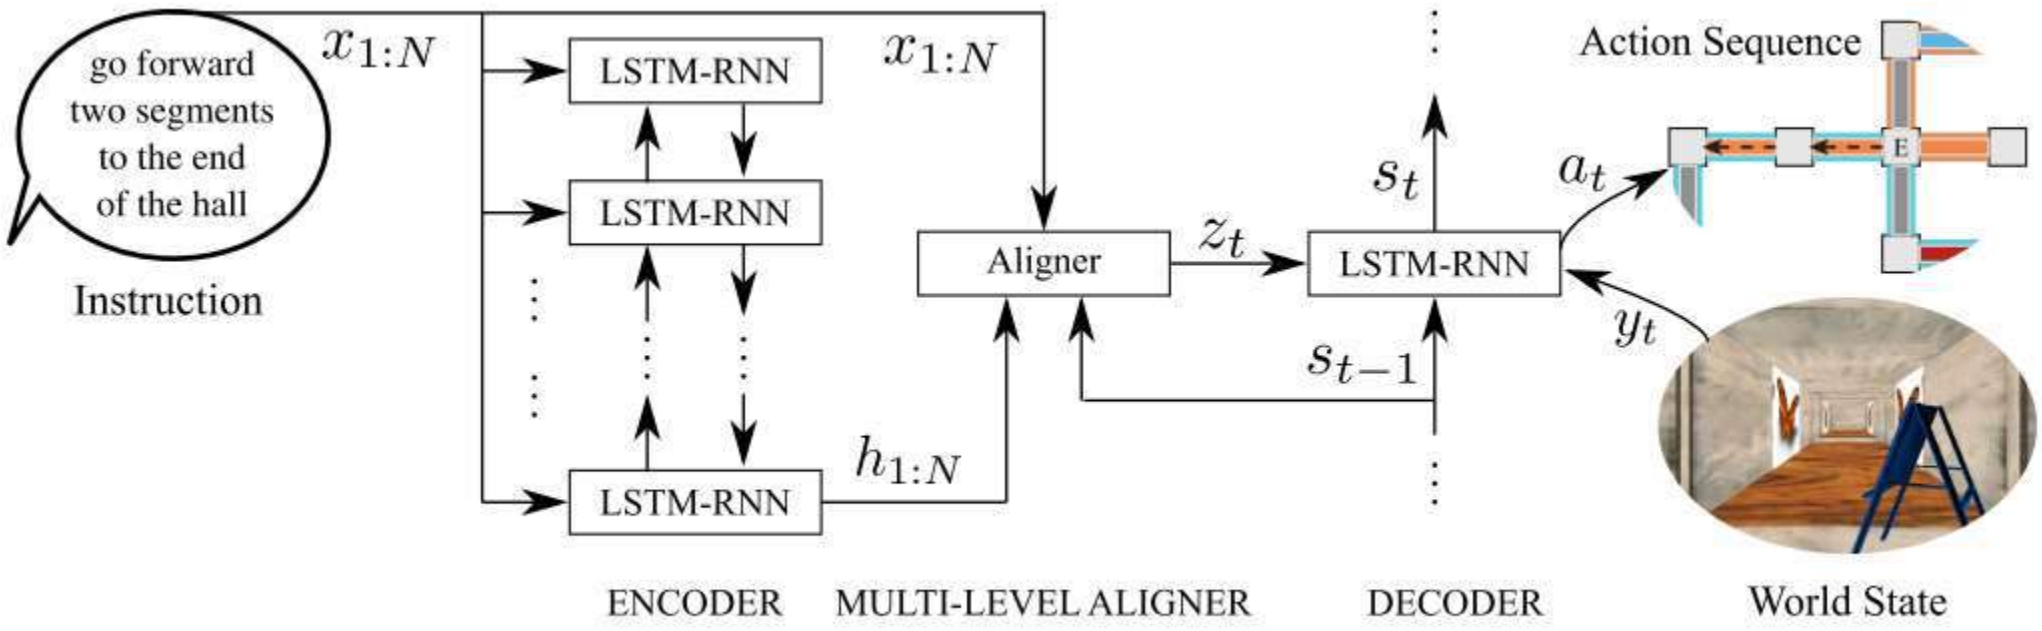

Figure 2: Our encoder-aligner-decoder model with multi-level alignment

Mei, et al. Listen, Attend, and Walk: Neural Mapping of Navigational Instructions to Action Sequences. AAAI 2016.

Look Before You Leap

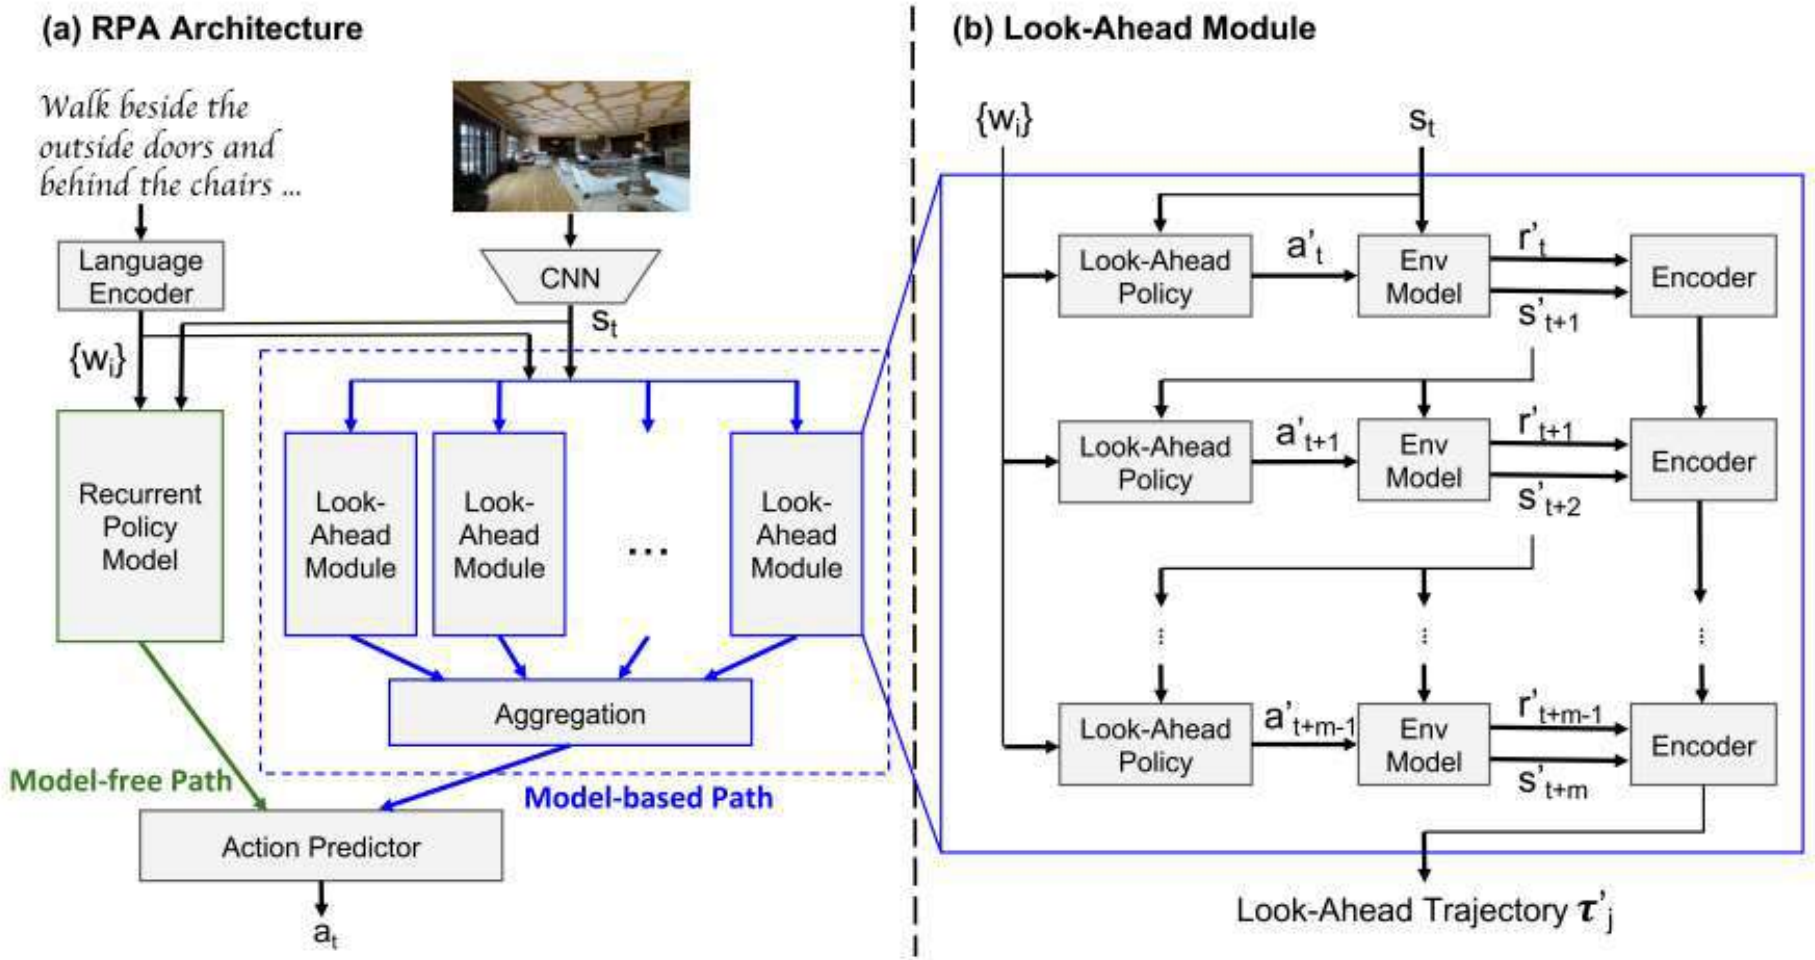

Wang, Xin, et al. Look Before You Leap: Bridging Model-Free and Model-Based Reinforcement Learning for Planned-Ahead Vision-and-Language Navigation. ECCV2018.

Look Before You Leap

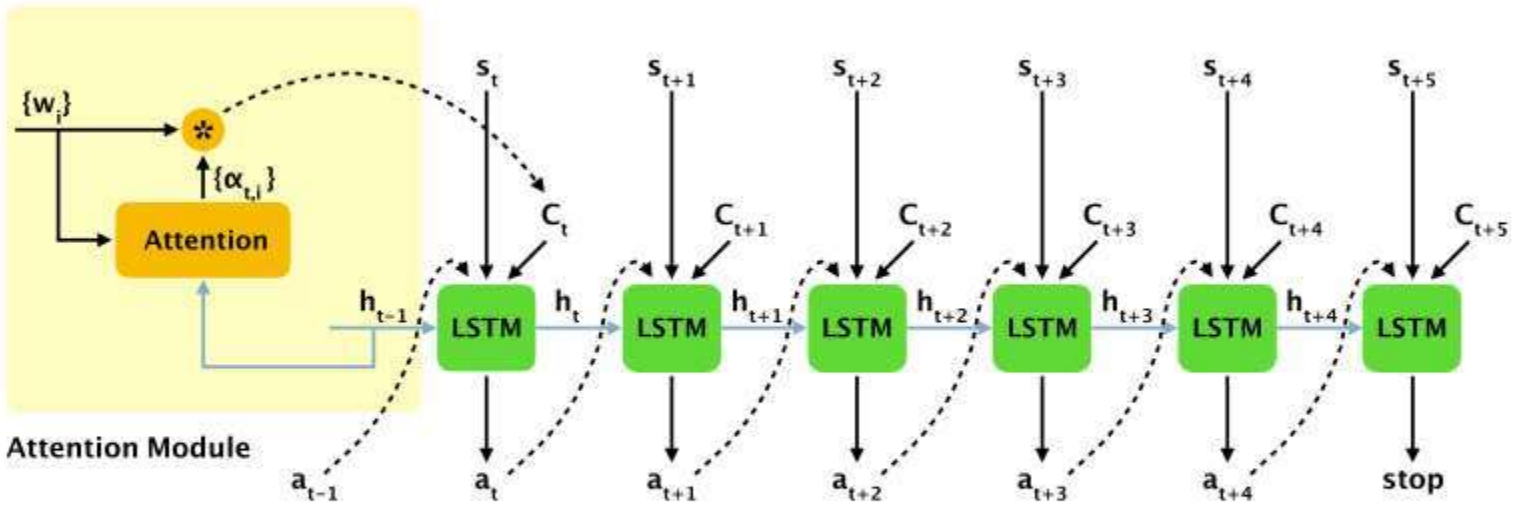

Fig. 4: An example of the unrolled recurrent policy model (from t to $t+5$). The left-side yellow region demonstrates the attention mechanism at time step t.

$$c_t = \sum \alpha_{t,i} w_i$$

$$\alpha_{t,i} = \frac{\exp(e_{t,i})}{\sum_{k=1}^{n} \exp(e_{t,k})} \quad , \quad \text{where } e_{t,i} = h_{t-1}^{\top} w_i$$

$$h_t = LSTM(h_{t-1}, [c_t, s_t, a_{t-1}])$$

Wang, Xin, et al. Look Before You Leap: Bridging Model-Free and Model-Based Reinforcement Learning for Planned-Ahead Vision-and-Language Navigation. ECCV2018.

Look Before You Leap

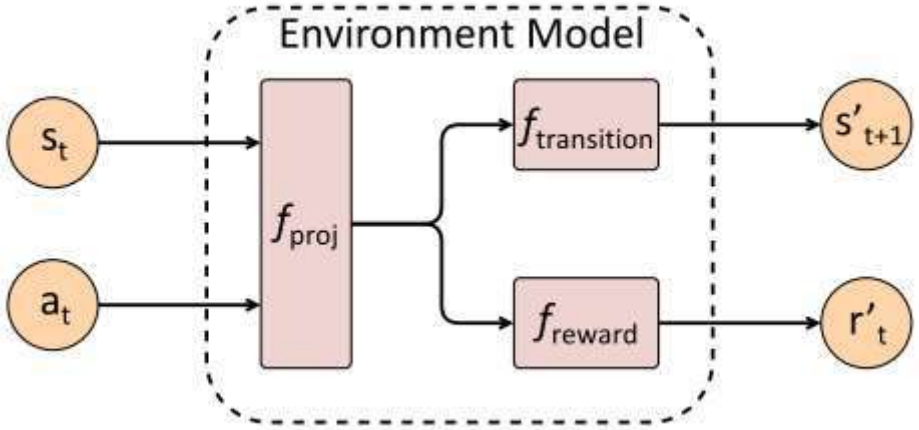

$$s'_{t+1} = f_{\text{transition}}(f_{\text{proj}}(s_t, a_t))$$
$$r'_t = f_{\text{reward}}(f_{\text{proj}}(s_t, a_t))$$

$$\max_{\theta} \mathcal{J}^{\pi} = \mathbb{E}\left[\sum_{t=1}^{T} \gamma^{t-1} r(a_t, s_t) | \pi(o_t; \theta)\right]$$

Wang, Xin, et al. Look Before You Leap: Bridging Model-Free and Model-Based Reinforcement Learning for Planned-Ahead Vision-and-Language Navigation. ECCV2018.

Speaker-Follower

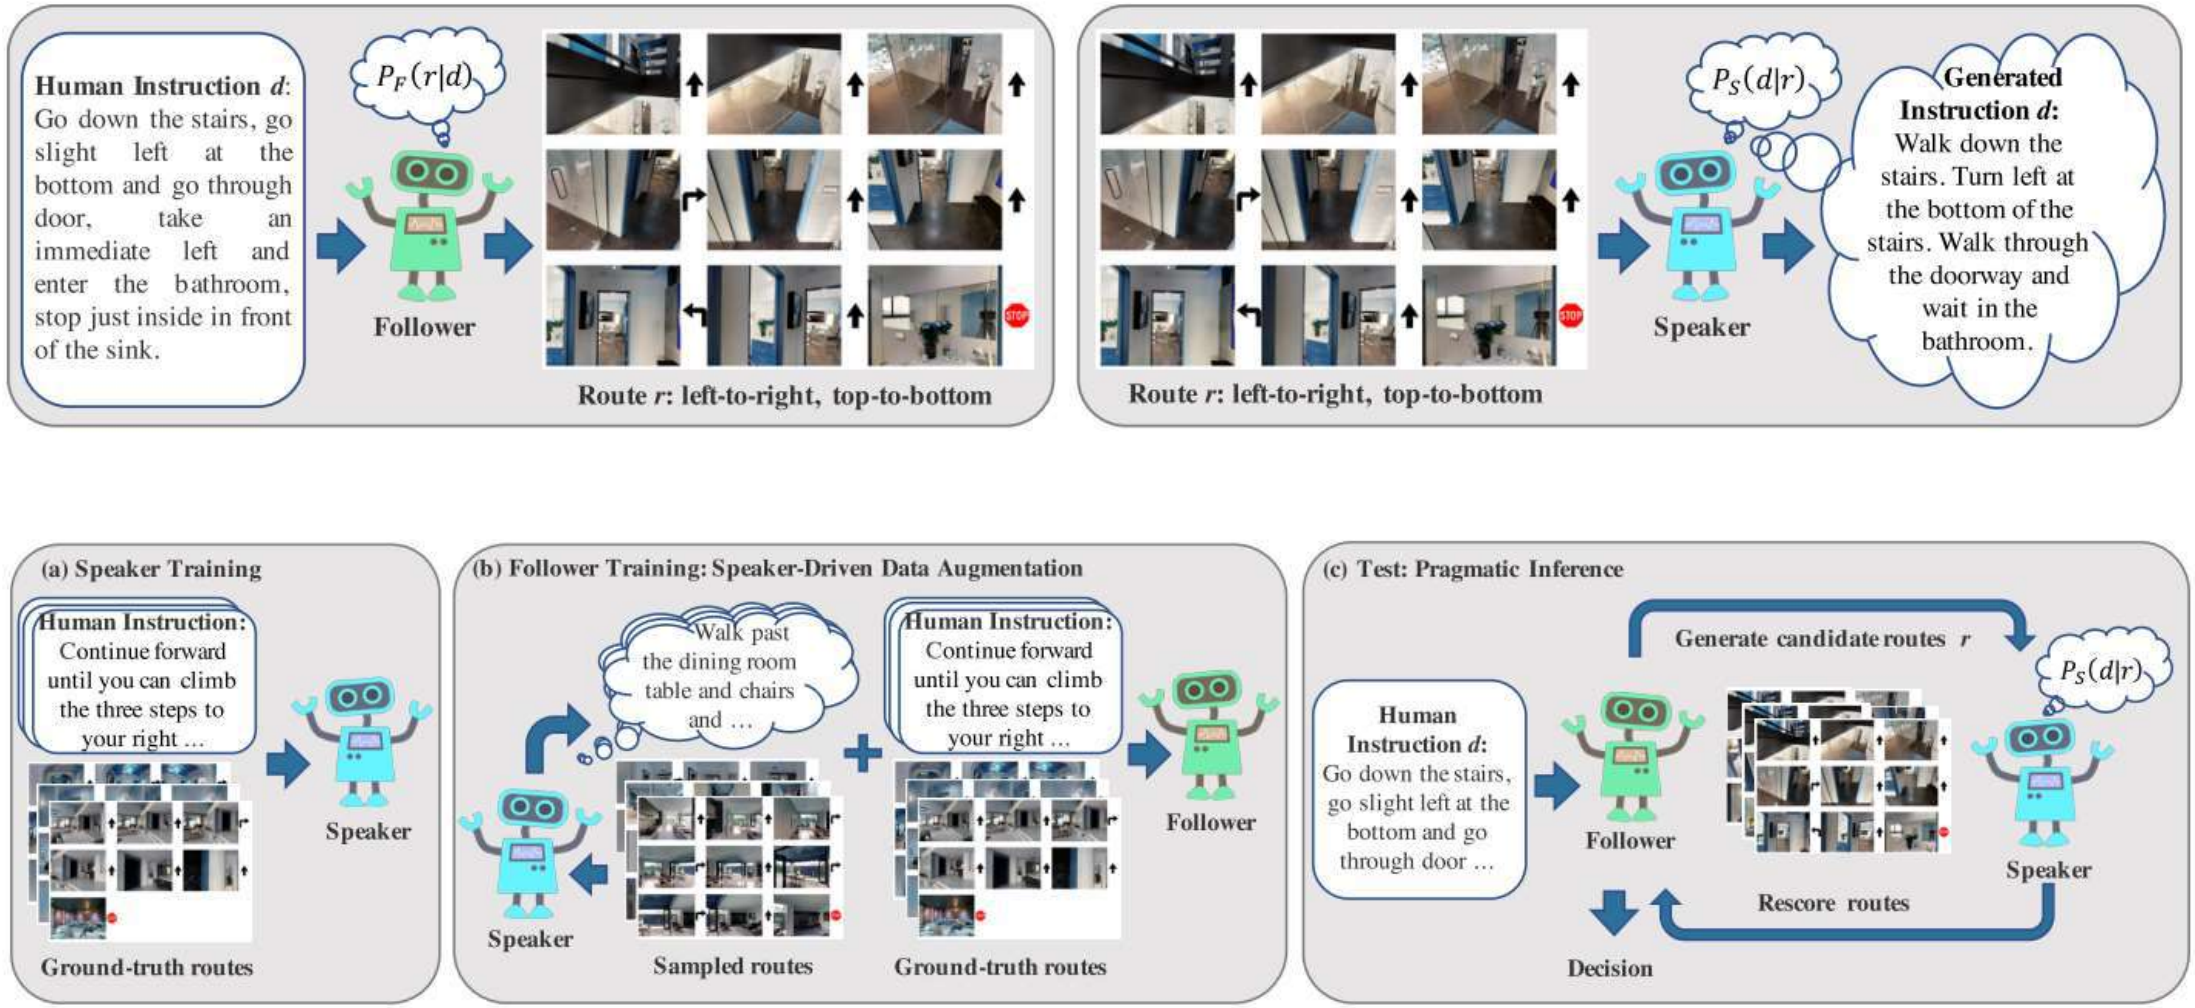

Fried, Daniel, et al. Speaker-Follower Models for Vision-and-Language Navigation. NIPS2018

Speaker-Follower

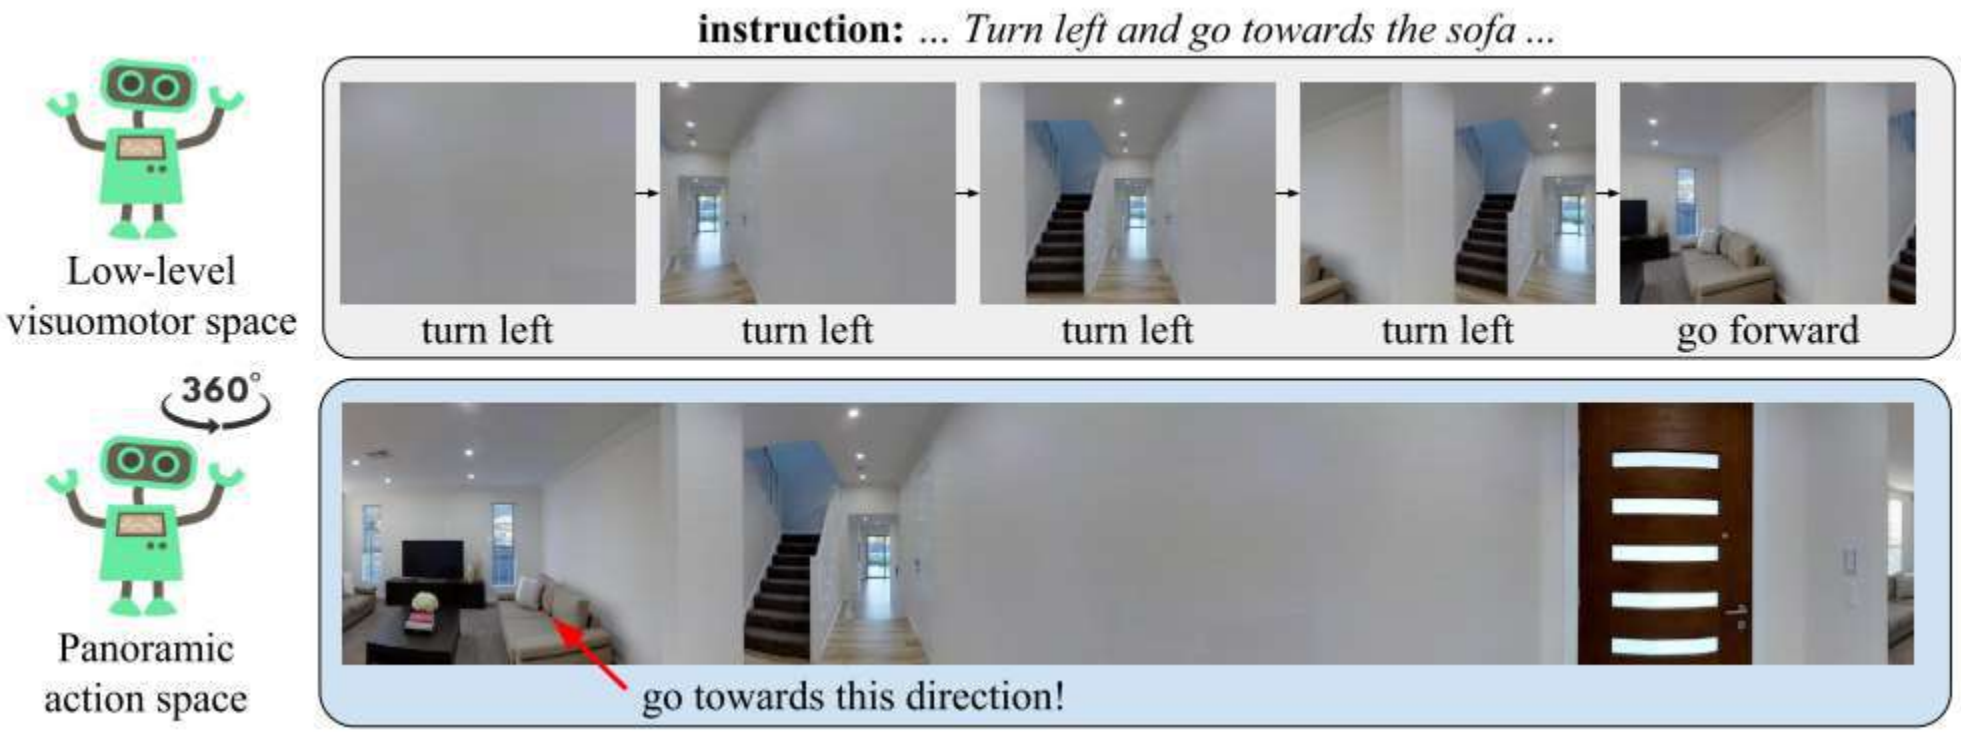

Reinforced Cross-Modal Match

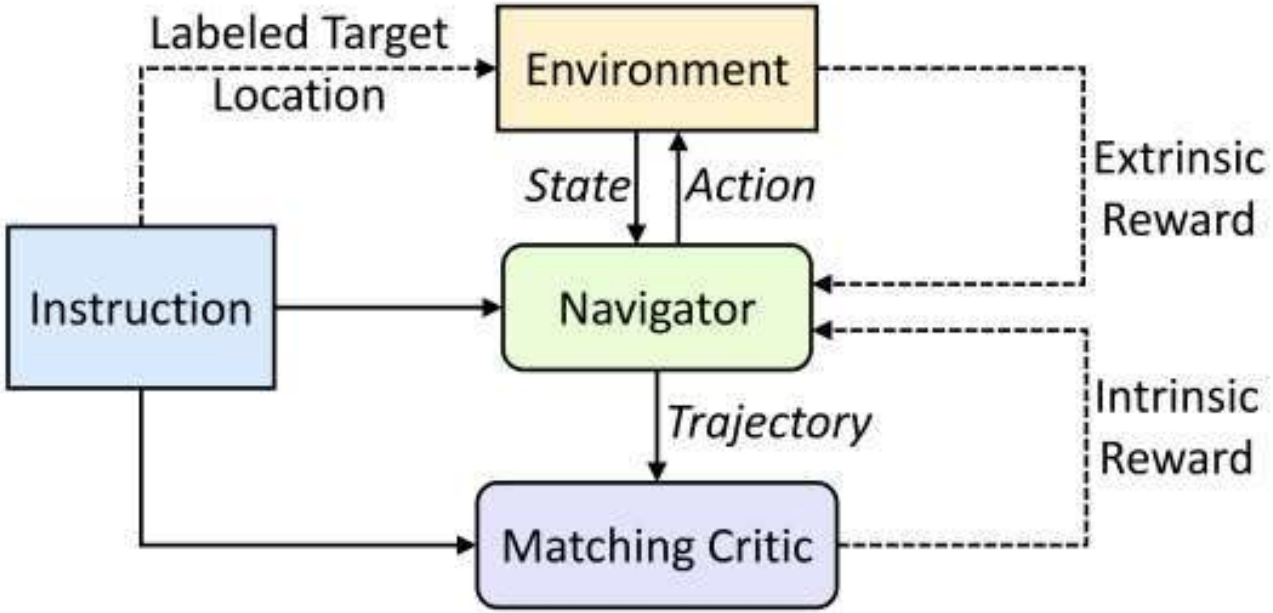

Figure 2: Overview of our RCM framework.

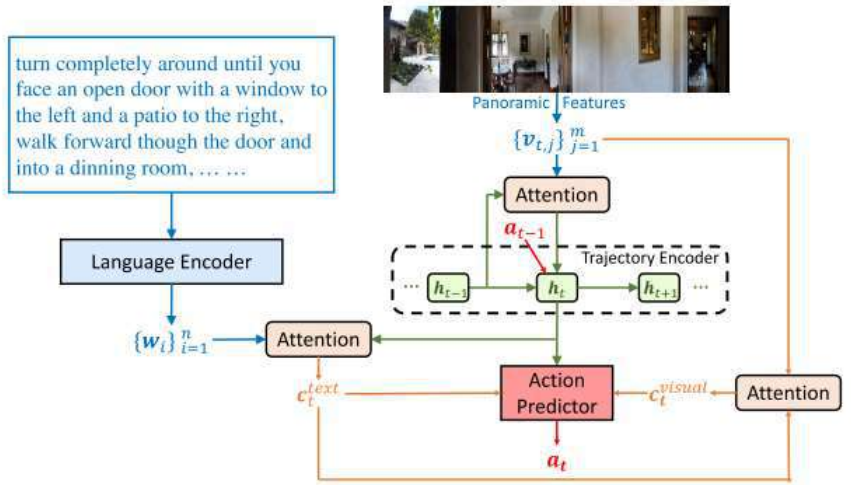

Figure 3: Cross-modal reasoning navigator at step t.

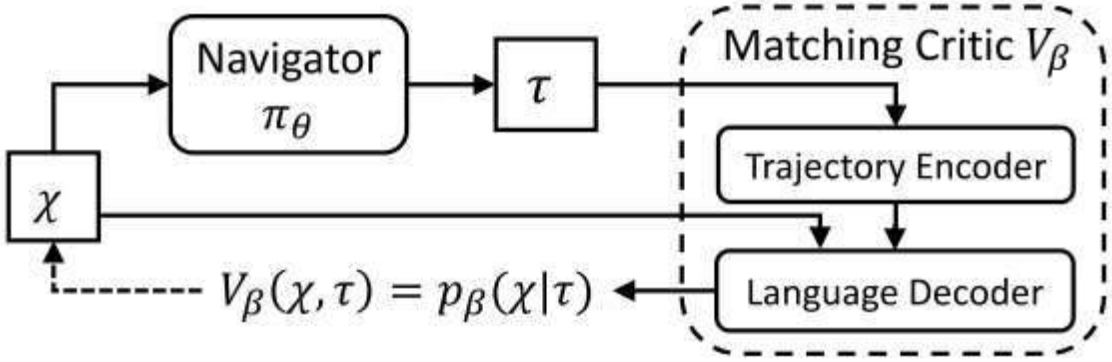

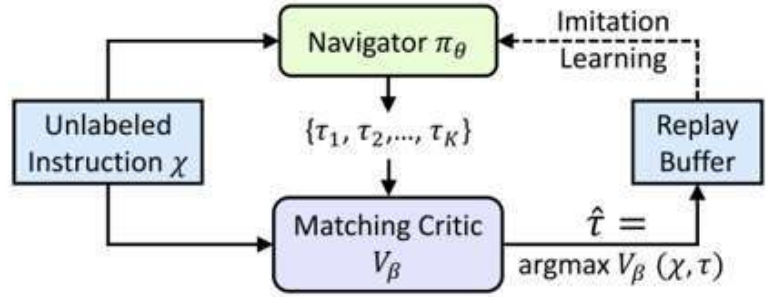

Wang Xin, et al. Reinforced Cross-Modal Matching and Self-Supervised Imitation Learning for Vision-Language Navigation CVPR2019 best stu paper

Self Monitoring

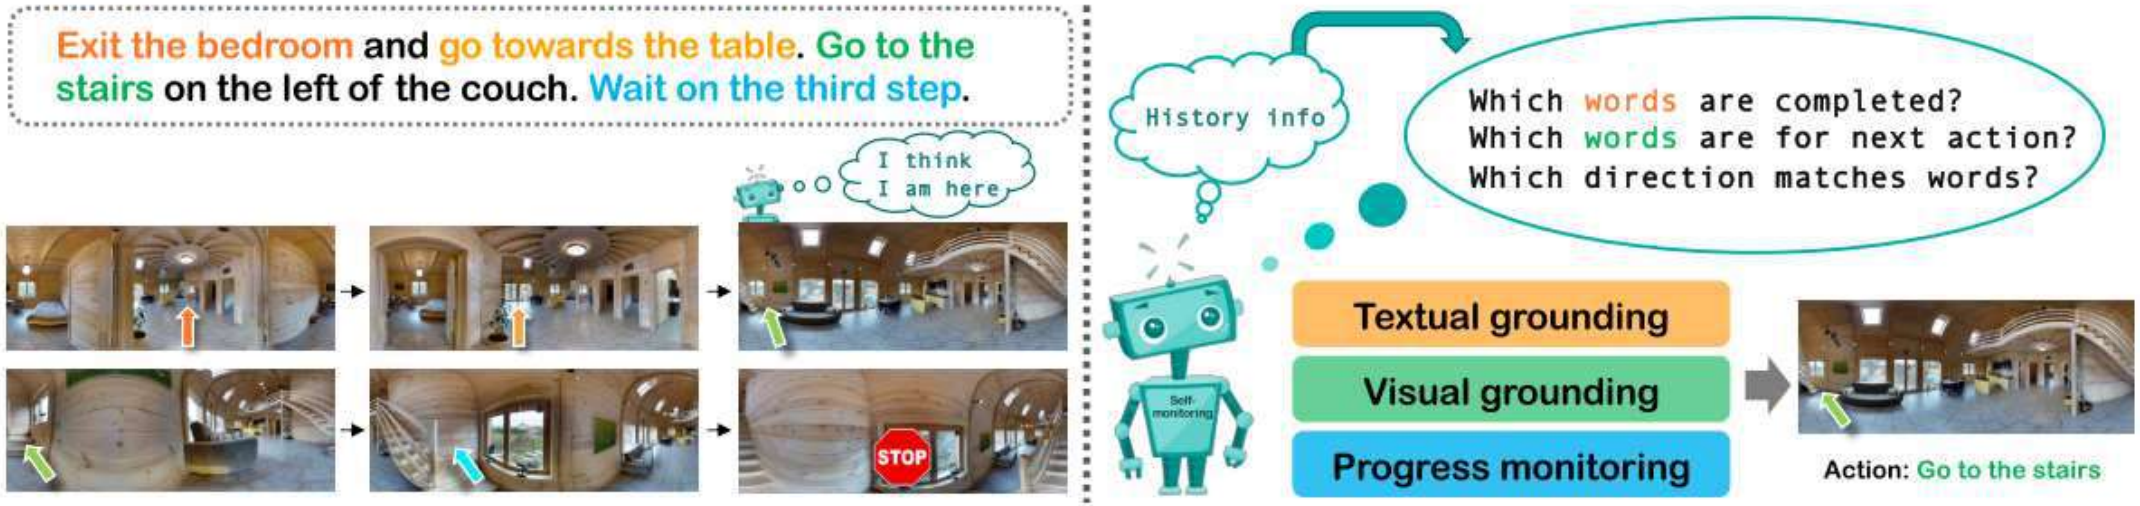

Ma, Chih-Yao, et al. SELF-MONITORING NAVIGATION AGENT VIA AUXIL-IARY PROGRESS ESTIMATION. ICLR2019.

Self Monitoring

$$z_{t,l}^{\text{textual}} = (\boldsymbol{W}_x \boldsymbol{h}_{t-1})^\top PE(\boldsymbol{x}_l) \qquad o_{t,k} = (\boldsymbol{W}_a[\boldsymbol{h}_t, \hat{\boldsymbol{x}}_t])^\top g(\boldsymbol{v}_{t,k}) \qquad z_{t,k}^{\text{visual}} = (\boldsymbol{W}_v \boldsymbol{h}_{t-1})^\top g(\boldsymbol{v}_{t,k})$$

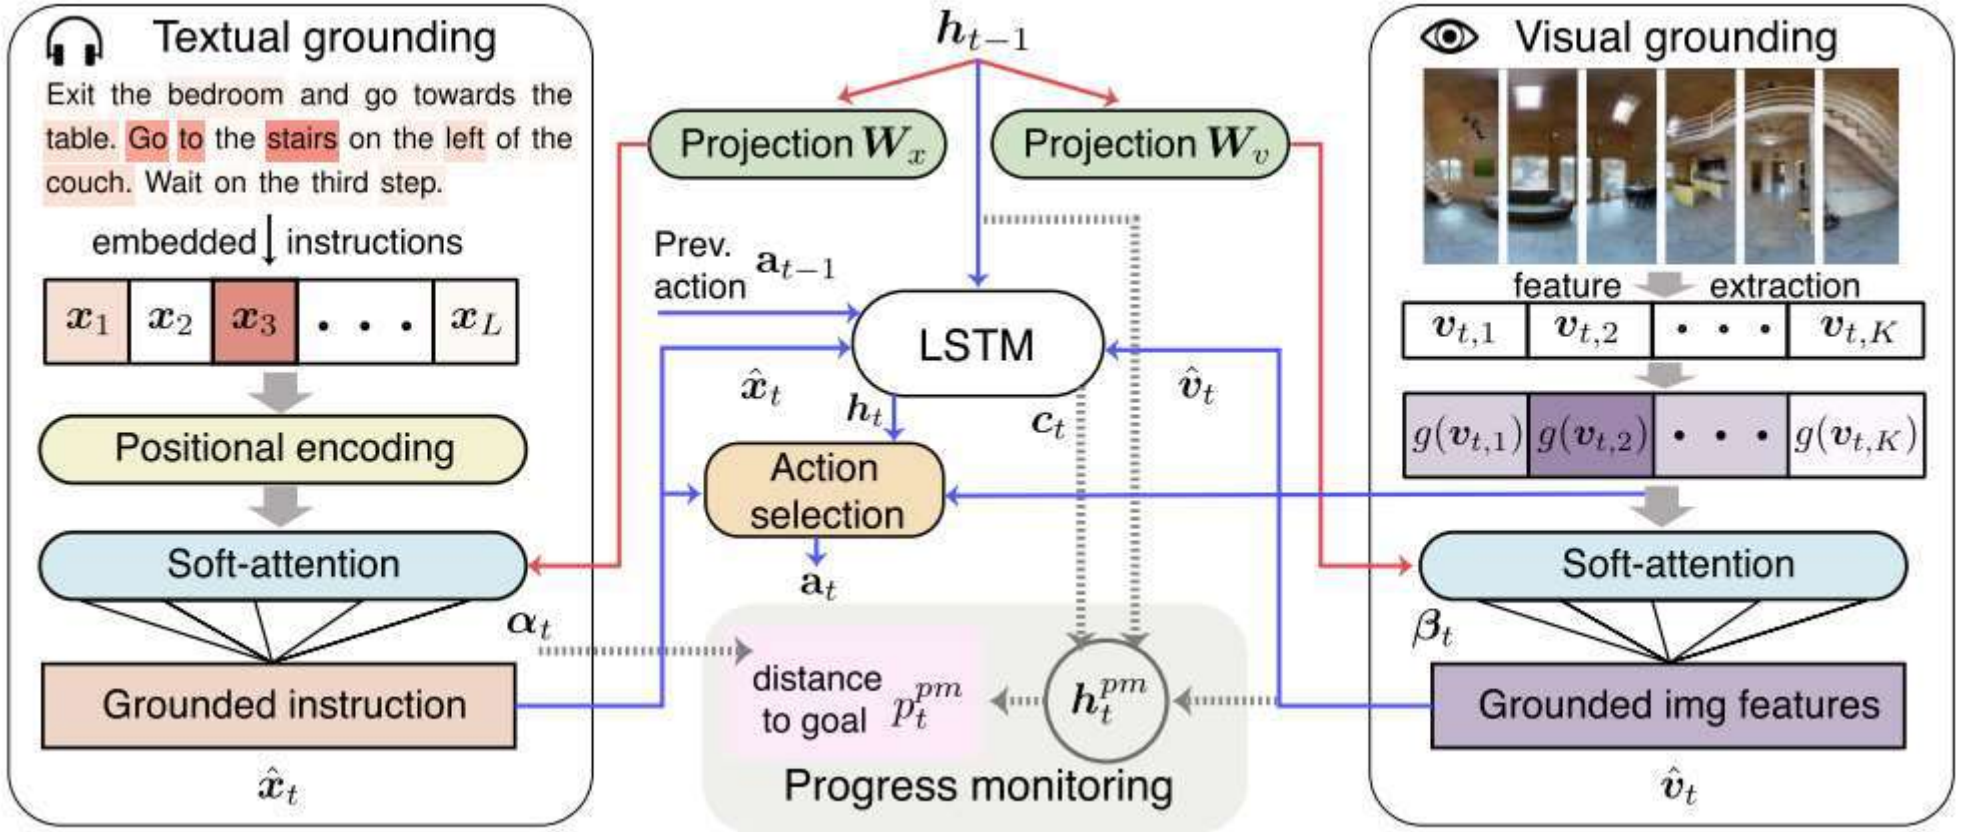

$$\boldsymbol{h}_t^{pm} = \sigma(\boldsymbol{W}_h([\boldsymbol{h}_{t-1}, \hat{\boldsymbol{v}}_t]) \otimes \tanh(\boldsymbol{c}_t))$$

Ma, Chih-Yao, et al. SELF-MONITORING NAVIGATION AGENT VIA AUXIL-IARY PROGRESS ESTIMATION. ICLR2019.

Environmental Dropout

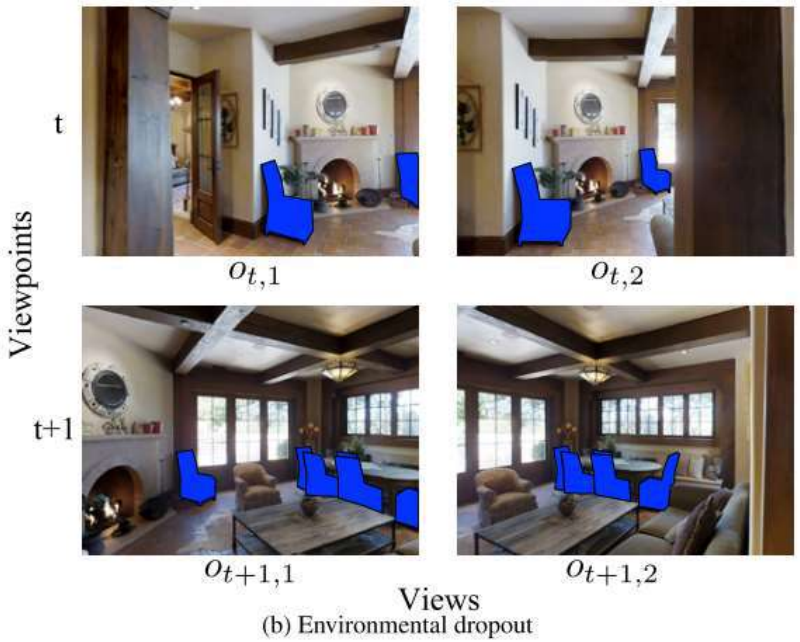

(b) Environmental dropout

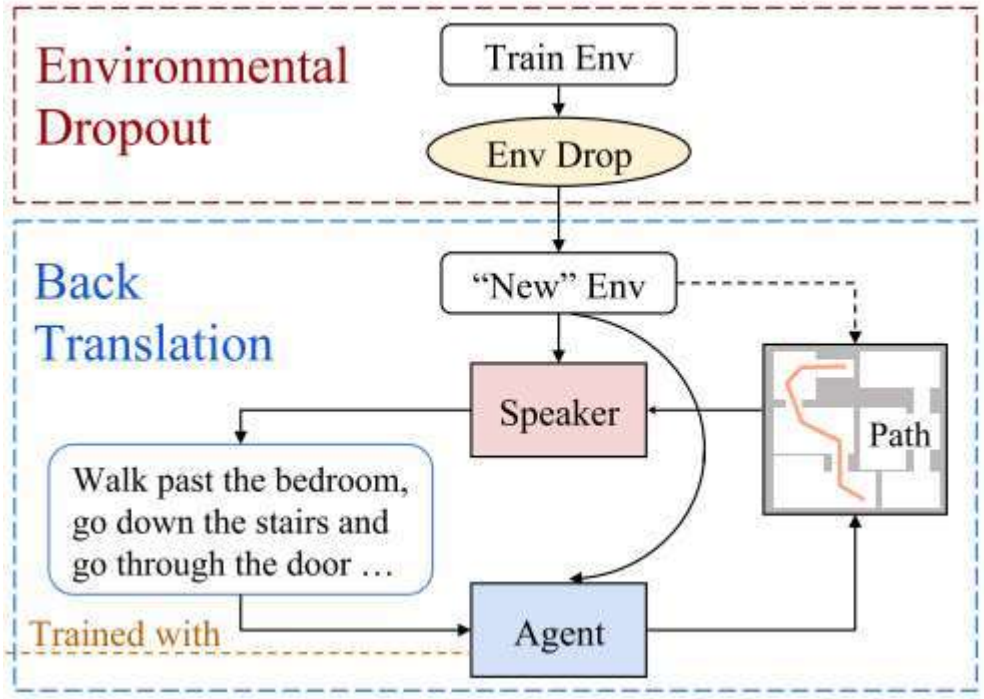

Tan Hao, et al. Learning to Navigate Unseen Environments: Back Translation with Environmental Dropout. ACL2019

Environmental Dropout

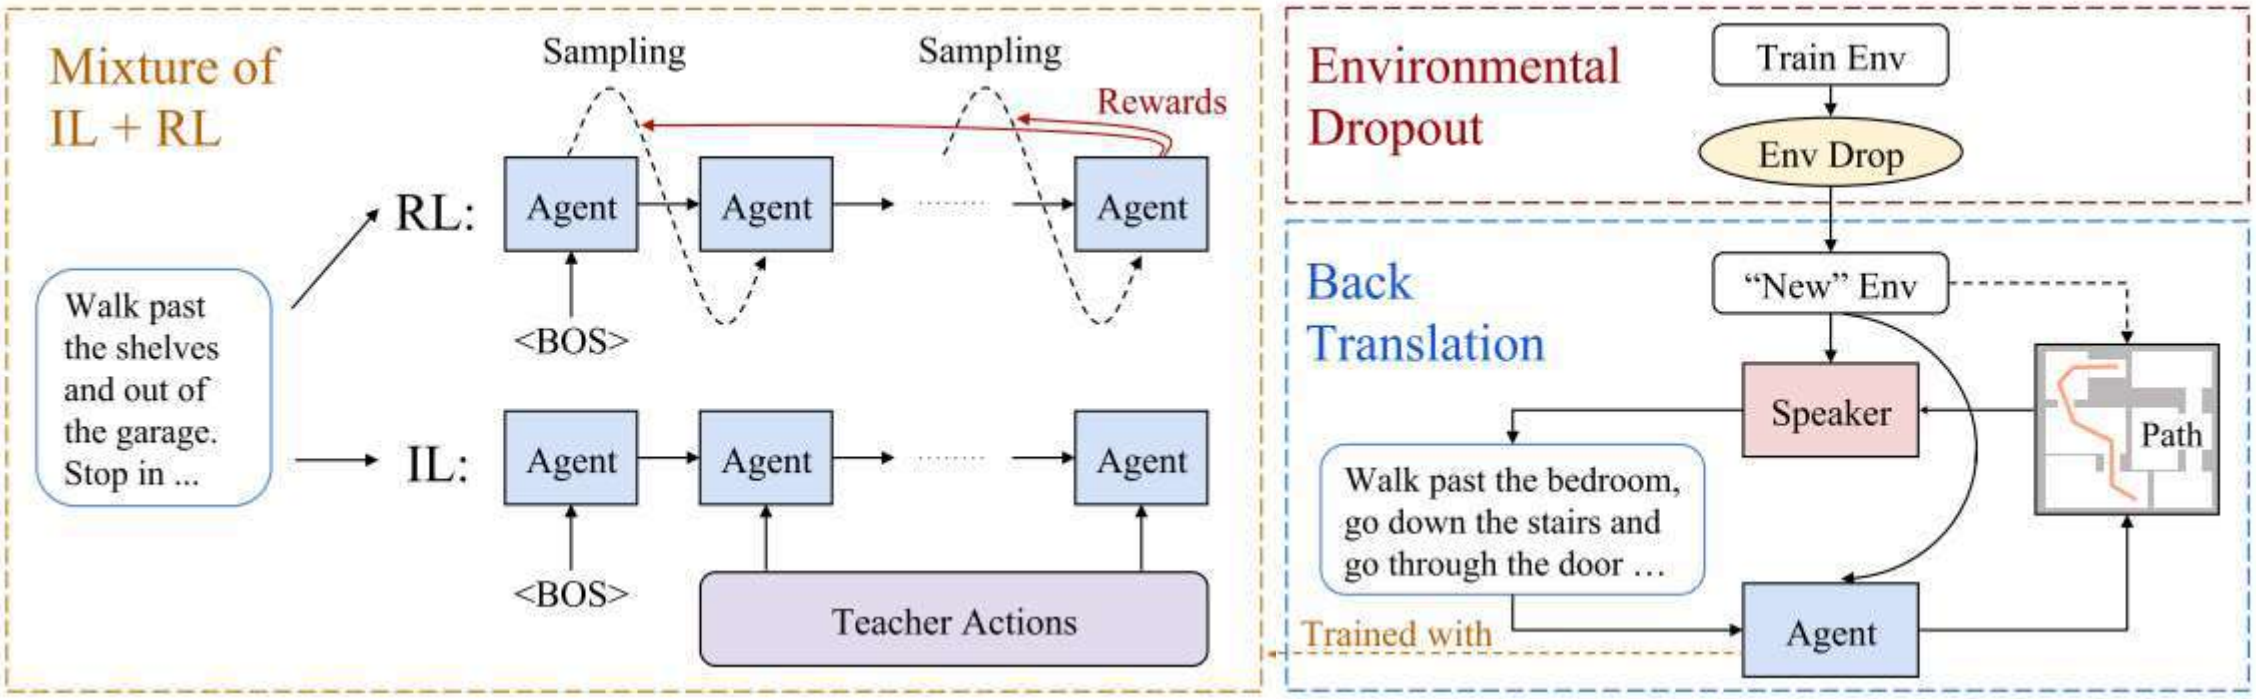

$$\mathcal{L}^{\text{MIX}} = \mathcal{L}^{\text{RL}} + \lambda_{\text{IL}} \mathcal{L}^{\text{IL}}$$

Tan Hao, et al. Learning to Navigate Unseen Environments: Back Translation with Environmental Dropout. ACL2019

Auxiliary Reasoning Tasks

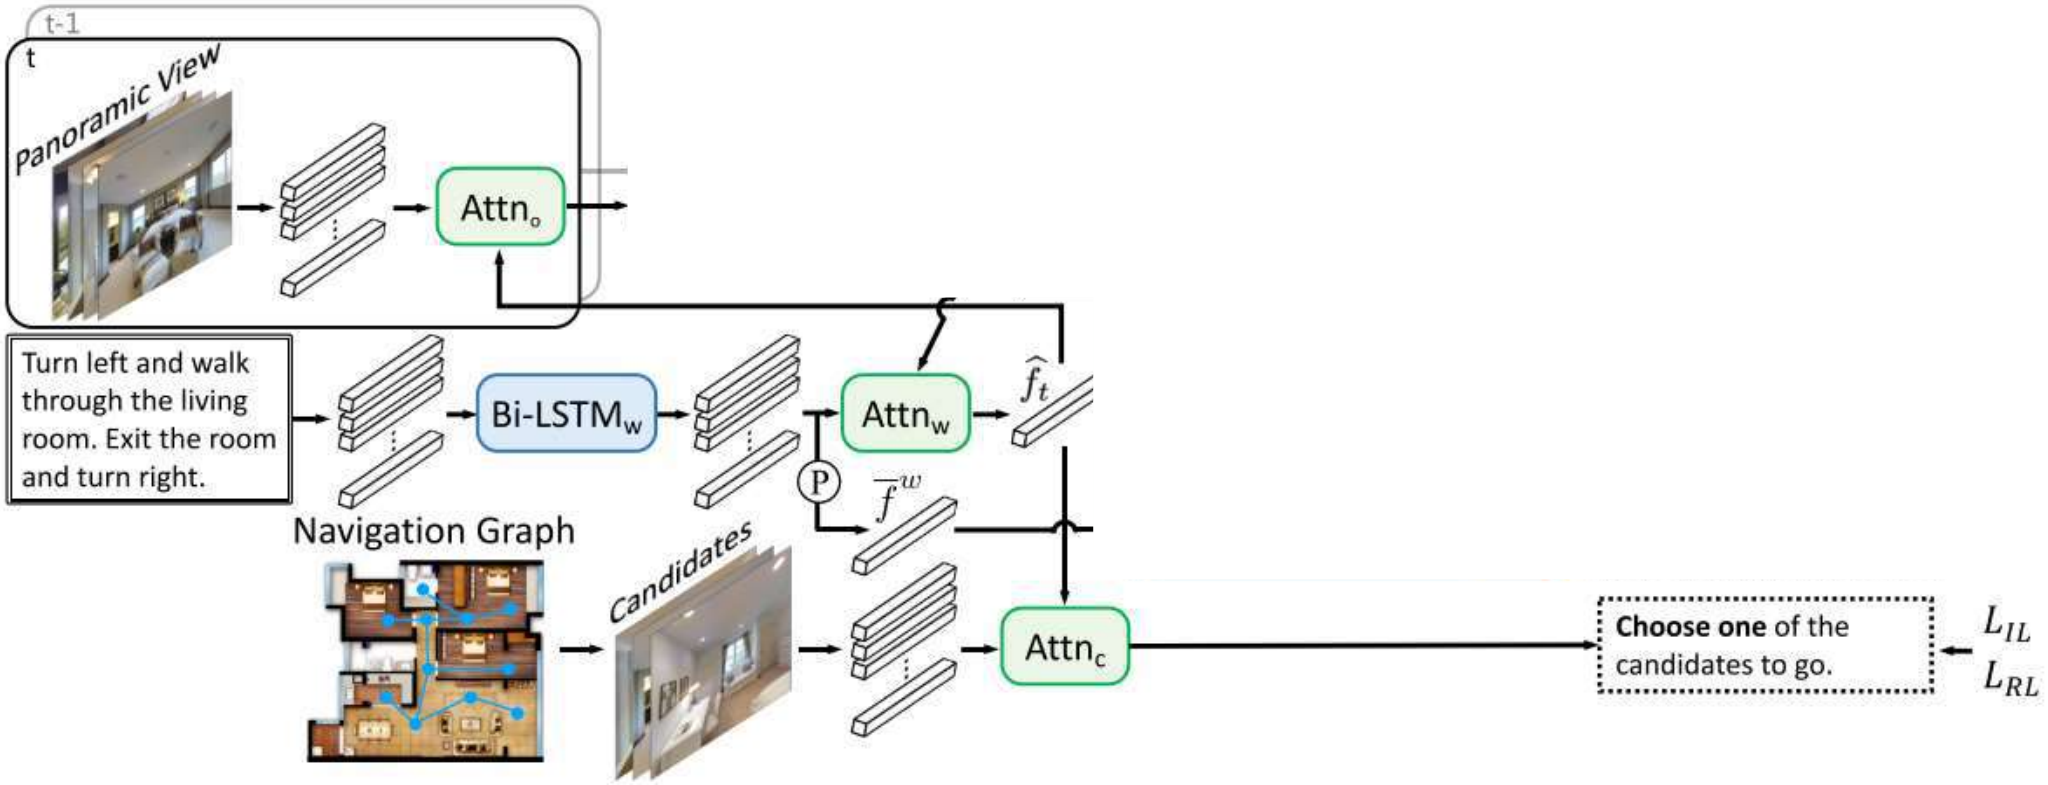

Zhu fd, Vision-Language Navigation with Self-Supervised Auxiliary Reasoning Tasks. CVPR2020

Auxiliary Reasoning Tasks

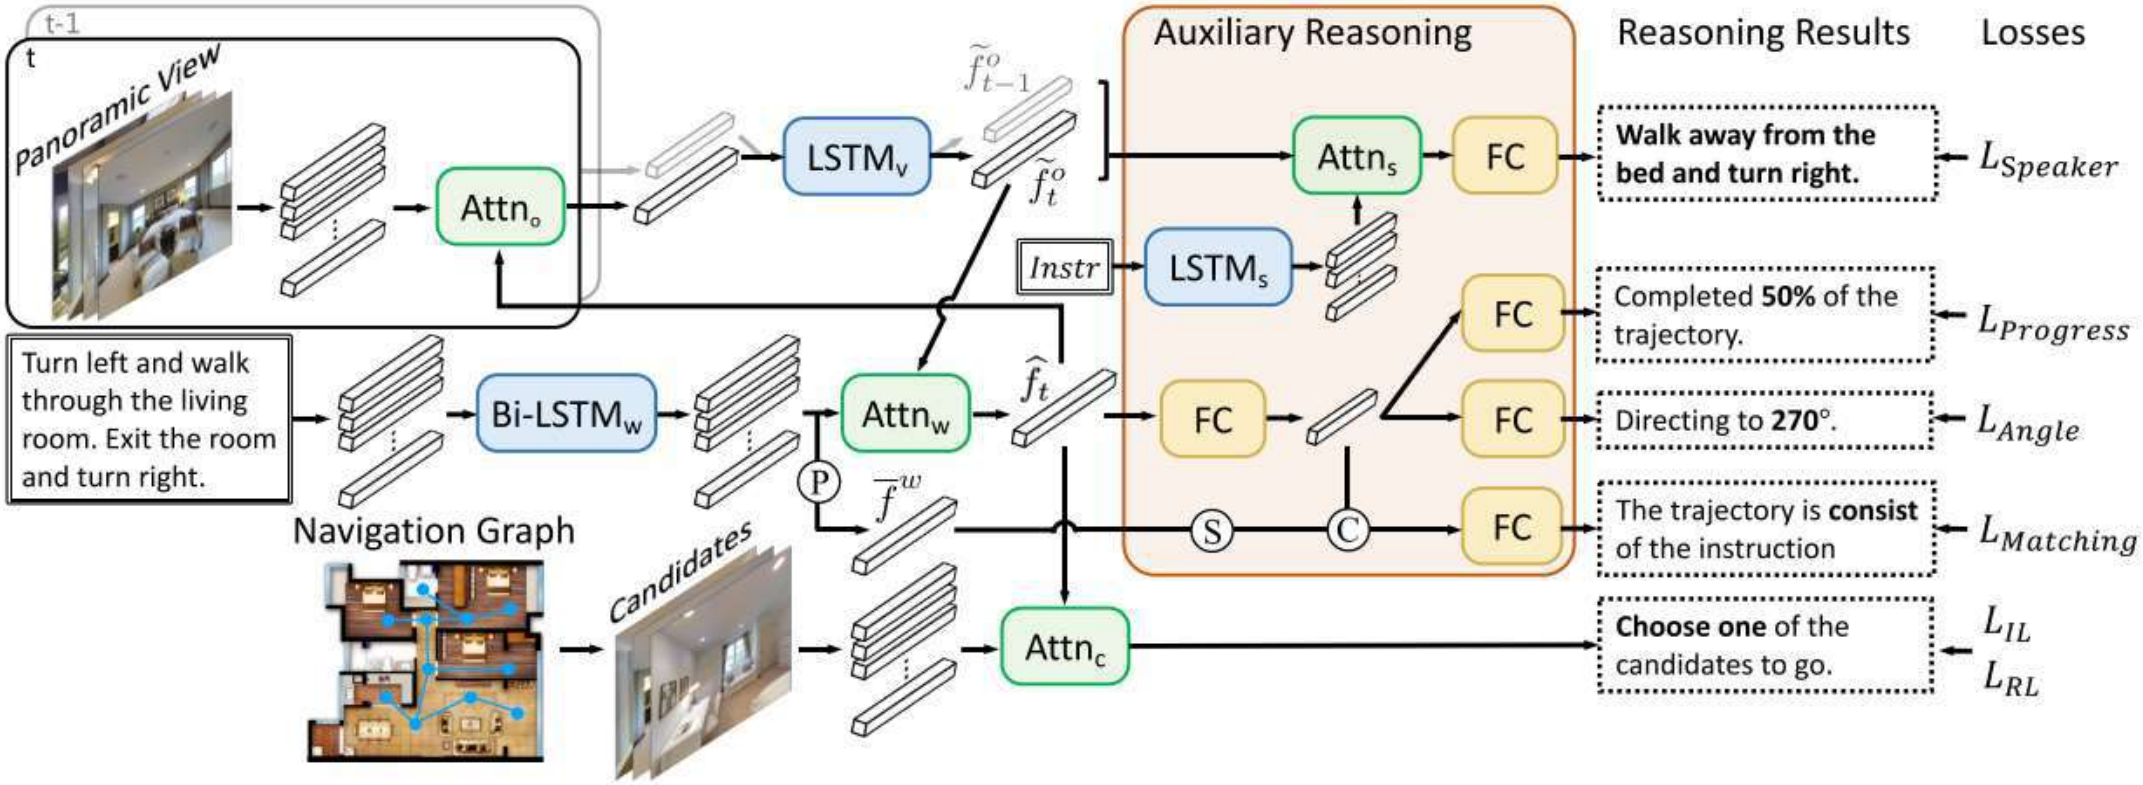

Zhu fd, Vision-Language Navigation with Self-Supervised Auxiliary Reasoning Tasks. CVPR2020

Auxiliary Reasoning Tasks

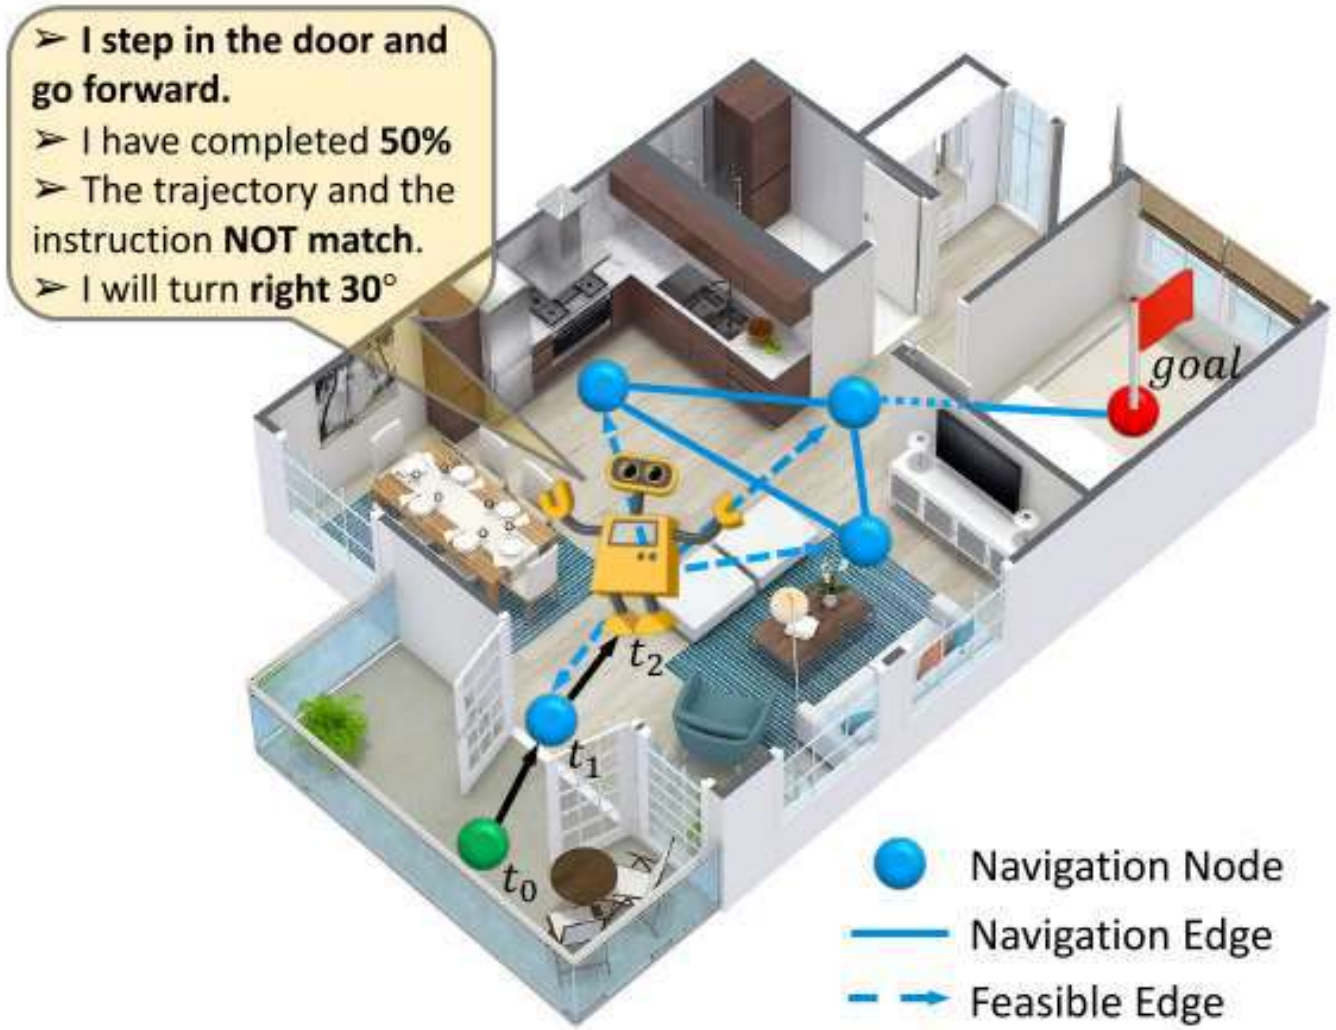

Zhu fd, Vision-Language Navigation with Self-Supervised Auxiliary Reasoning Tasks. CVPR2020

Experimental results

Leader-Board (Test Unseen)	Single Run			
Models	NE	OR	SR	*SPL*
Random [5]	9.79	0.18	0.17	0.12
Seq-to-Seq [5]	20.4	0.27	0.20	0.18
Look Before You Leap [42]	7.5	0.32	0.25	0.23
Speaker-Follower [10]	6.62	0.44	0.35	0.28
Self-Monitoring [23]	5.67	0.59	0.48	0.35
Reinforced Cross-Modal [41]	6.12	0.50	0.43	0.38
Environmental Dropout [37]	5.23	0.59	0.51	0.47
AuxRN(Ours)	**5.15**	**0.62**	**0.55**	**0.51**

Zhu fd, Vision-Language Navigation with Self-Supervised Auxiliary Reasoning Tasks. CVPR2020

Conclusion

- Base: Seq-to-seq

- Align(Matching, Progress)
- Data augmentation(Self/semi-supervise)
- RL(Model-based/free, IL, SIL)

- How to get the data

Thank you

- End